

Exercise 1

a) The pinball loss is :

$$L(y, f(x)) = \begin{cases} -(1-\alpha)(y-f(x)) & , y < f(x) \\ \alpha(y-f(x)) & , y \geq f(x) \end{cases}$$

$$g_{1(i)} = -\frac{\partial L(y_i, f(x_i))}{\partial f(x_i)} = \begin{cases} \frac{\partial (1-\alpha)f(x)}{\partial f(x)} \\ \frac{\partial (-\alpha f(x))}{\partial f(x)} \end{cases}$$

$$\Rightarrow \begin{cases} 1-\alpha & , y_i < f(x_i) \\ -\alpha & , y_i \geq f(x_i) \end{cases}$$

b) Let, $f(x_i) = \theta_0 + \theta_1 x_i$
 g_i is defined as :

$$-g_i = \frac{\partial L}{\partial \theta_0} ; -g_i x_i = \frac{\partial L}{\partial \theta_1}$$

The updated rules are \rightarrow

$$\theta_0 \rightarrow \theta_0 - \eta \frac{1}{n} \sum_{i=1}^n (-g_i) = \theta_0 + \eta \frac{1}{n} \sum_{i=1}^n g_i$$

$$\theta_1 \rightarrow \theta_1 - \eta \frac{1}{n} \sum_{i=1}^n (-g_i x_i) = \theta_1 + \eta \frac{1}{n} \sum_{i=1}^n g_i x_i$$

Homework 6

C

```
data <- read.csv("gd.csv")

theta_0 <- 2
theta_1 <- 2
alpha <- 0.7
eta <- 0.8
n_iter <- 25
n <- nrow(data)

theta_0_values <- numeric(n_iter + 1)
theta_1_values <- numeric(n_iter + 1)
empirical_risk <- numeric(n_iter + 1)

#pinball loss
pinball_loss <- function(y, fx, alpha) {
  if (y < fx) {
    return((1 - alpha) * (fx - y))
  } else {
    return(alpha * (y - fx))
  }
}

#j = 0
theta_0_values[1] <- theta_0
theta_1_values[1] <- theta_1

loss_sum <- 0
for (i in 1:n) {
  x_i <- data$x[i]
  y_i <- data$y[i]
  fx <- theta_0 + theta_1 * x_i
  loss_sum <- loss_sum + pinball_loss(y_i, fx, alpha)
}
empirical_risk[1] <- loss_sum

# Gradient descent
for (j in 1:n_iter) {
  for (i in 1:n) {
    x_i <- data$x[i]
    y_i <- data$y[i]
    fx <- theta_0 + theta_1 * x_i

    grad <- if (y_i < fx) { 1 - alpha } else { -alpha }

    theta_0 <- theta_0 - eta * grad
```

```

    theta_1 <- theta_1 - eta * grad * x_i
  }

  theta_0_values[j + 1] <- theta_0
  theta_1_values[j + 1] <- theta_1

  loss_sum <- 0
  for (i in 1:n) {
    x_i <- data$x[i]
    y_i <- data$y[i]
    fx <- theta_0 + theta_1 * x_i
    loss_sum <- loss_sum + pinball_loss(y_i, fx, alpha)
  }
  empirical_risk[j + 1] <- loss_sum
}

results <- data.frame(
  Iteration = 0:n_iter,
  theta_0 = theta_0_values,
  theta_1 = theta_1_values,
  Empirical_Risk = empirical_risk
)

print(results)

```

	Iteration	theta_0	theta_1	Empirical_Risk
1	0	2.0	2.000000	113.60258
2	1	2.8	5.306699	41.73446
3	2	2.8	5.146467	40.98632
4	3	2.8	5.396298	42.40155
5	4	2.8	5.304486	41.71798
6	5	2.8	5.144254	40.97840
7	6	2.8	5.394084	42.38507
8	7	2.8	5.302272	41.70150
9	8	2.8	5.142040	40.97047
10	9	2.8	5.307286	41.73883
11	10	2.8	5.147054	40.98843
12	11	2.8	5.396884	42.40592
13	12	2.8	5.305072	41.72235
14	13	2.8	5.144841	40.98050
15	14	2.8	5.394671	42.38944
16	15	2.8	5.302859	41.70587
17	16	2.8	5.142627	40.97257
18	17	2.8	5.307873	41.74320
19	18	2.8	5.147641	40.99053
20	19	2.8	5.397471	42.41029
21	20	2.8	5.305659	41.72672
22	21	2.8	5.145427	40.98260
23	22	2.8	5.395258	42.39381
24	23	2.8	5.303446	41.71024


25	24	2.8	5.143214	40.97467
26	25	2.8	5.393045	42.37733

D

```
library(ggplot2)

ggplot(data, aes(x = x, y = y)) +
  geom_point(color = "black", alpha = 0.7) +
  lapply(0:n_iter, function(j) {
    theta0 <- theta_0_values[j + 1]
    theta1 <- theta_1_values[j + 1]
    color <- rgb(j / n_iter, 0, 1 - j / n_iter)
    geom_abline(intercept = theta0, slope = theta1, color = color, size = 0.8,
alpha = 0.7)
  }) +
  labs(title = "Evolution of the Regression Line during Gradient Descent",
       x = "x", y = "y") +
  theme_minimal()
```

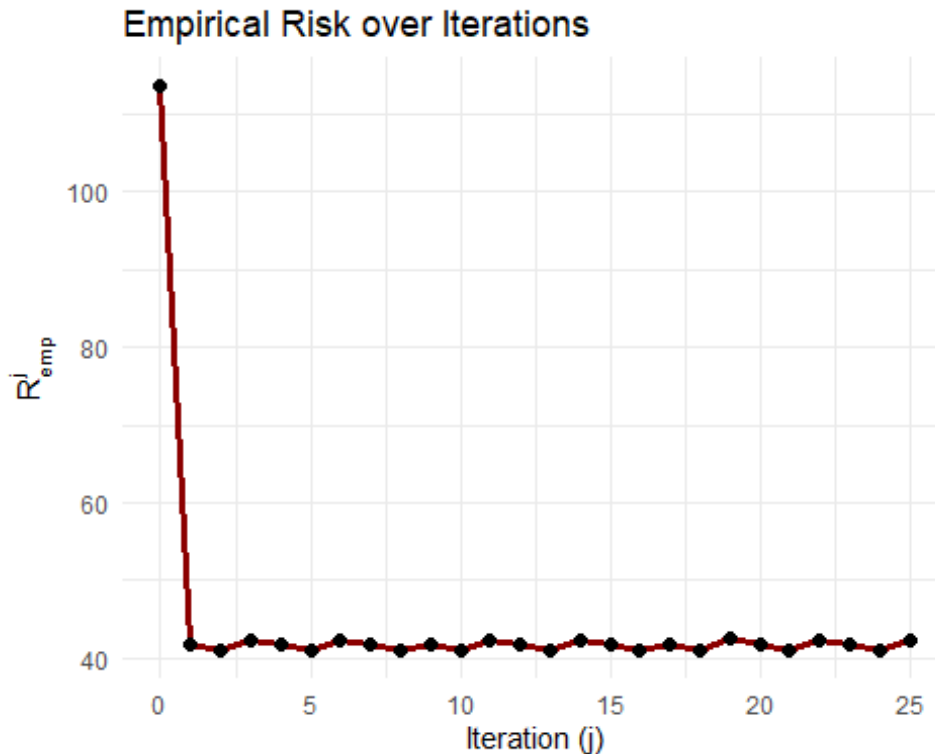
Warning: Using `size` aesthetic for lines was deprecated in ggplot2 3.4.0.
 i Please use `linewidth` instead.



E

```
library(ggplot2)
```

```
ggplot(results, aes(x = Iteration, y = Empirical_Risk)) +
  geom_line(color = "darkred", linewidth = 1.2) +
  geom_point(color = "black", size = 2) +
  labs(title = "Empirical Risk over Iterations",
       x = "Iteration (j)",
       y = expression(R[emp]^j)) +
  theme_minimal()
```



F

```
alpha <- 0.3
eta <- 0.8
n_iter <- 25
n <- nrow(data)

theta_0 <- 2
theta_1 <- 2

theta_0_values_03 <- numeric(n_iter + 1)
theta_1_values_03 <- numeric(n_iter + 1)
empirical_risk_03 <- numeric(n_iter + 1)

pinball_loss <- function(y, fx, alpha) {
  if (y < fx) {
    return((1 - alpha) * (fx - y))
  } else {
    return(alpha * (y - fx))
  }
}
```

```

}
}

theta_0_values_03[1] <- theta_0
theta_1_values_03[1] <- theta_1

loss_sum <- 0
for (i in 1:n) {
  x_i <- data$x[i]
  y_i <- data$y[i]
  fx <- theta_0 + theta_1 * x_i
  loss_sum <- loss_sum + pinball_loss(y_i, fx, alpha)
}
empirical_risk_03[1] <- loss_sum


for (j in 1:n_iter) {
  for (i in 1:n) {
    x_i <- data$x[i]
    y_i <- data$y[i]
    fx <- theta_0 + theta_1 * x_i
    grad <- if (y_i < fx) { 1 - alpha } else { -alpha }
    theta_0 <- theta_0 - eta * grad
    theta_1 <- theta_1 - eta * grad * x_i
  }

  theta_0_values_03[j + 1] <- theta_0
  theta_1_values_03[j + 1] <- theta_1

  loss_sum <- 0
  for (i in 1:n) {
    x_i <- data$x[i]
    y_i <- data$y[i]
    fx <- theta_0 + theta_1 * x_i
    loss_sum <- loss_sum + pinball_loss(y_i, fx, alpha)
  }
  empirical_risk_03[j + 1] <- loss_sum
}

ggplot(data, aes(x = x, y = y)) +
  geom_point(color = "black", alpha = 0.7) +
  lapply(0:n_iter, function(j) {
    theta0 <- theta_0_values_03[j + 1]
    theta1 <- theta_1_values_03[j + 1]
    color <- rgb(j / n_iter, 0, 1 - j / n_iter)
    geom_abline(intercept = theta0, slope = theta1, color = color, linewidth =
0.8, alpha = 0.7)
  }) +
  labs(title = "Evolution of the Regression Line during Gradient Descent ( $\alpha =$ 
0.3)",
       x = "x", y = "y") +
  theme_minimal()

```



When repeating the estimation with $\alpha = 0.3$, the resulting regression line after 25 iterations differs significantly from the one obtained with $\alpha = 0.7$.

Since $\alpha = 0.3$ penalizes over-predictions more heavily ($y < f(x)$), the gradient descent updates aim to reduce the model's tendency to overshoot. As a result, the final model has a **lower slope** and lies **below** the $\alpha = 0.7$ line.

In the graph, we observe that the blue line ($\alpha = 0.3$) is flatter and closer to the lower quantiles of the data, whereas the red line ($\alpha = 0.7$) follows a steeper trend, capturing higher quantiles.

This behavior is expected, as the quantile loss shifts the fitted model depending on α : smaller α values lead to more conservative models, favoring lower quantiles.

2

```
data <- read.csv("pw.csv")

fit_piecewise_poly <- function(data, K, degree) {
  breaks <- seq(0, 10, length.out = K + 1)
  data$interval <- cut(data$x, breaks = breaks, include.lowest = TRUE, right = FALSE)

  models <- list()
  fitted_data <- data

  for (k in levels(data$interval)) {
    subset_k <- subset(data, interval == k)
```

```

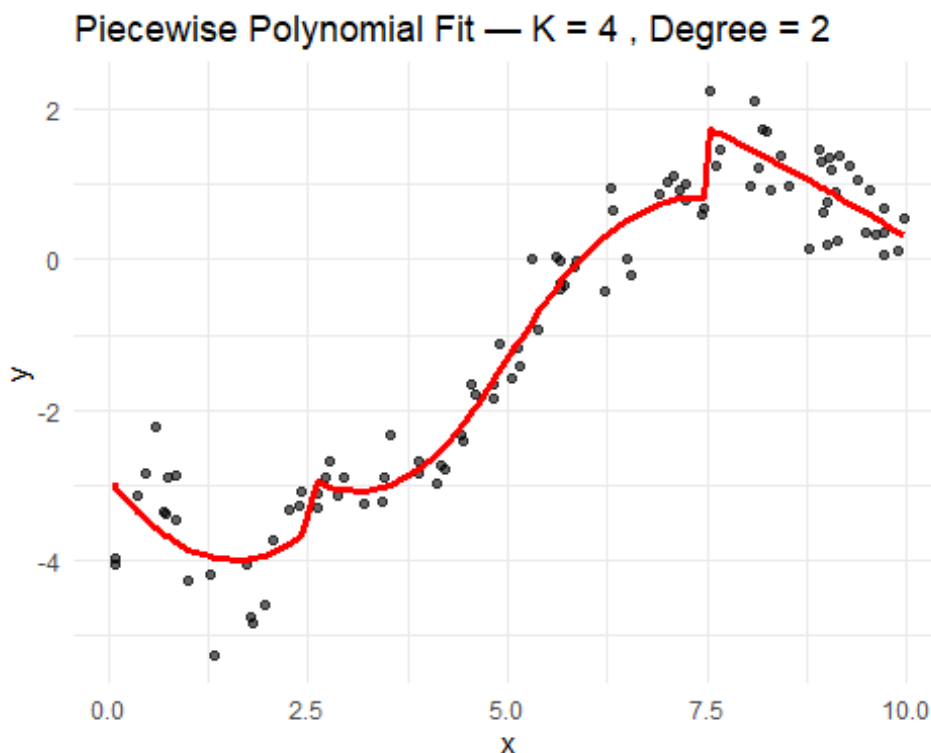
if (nrow(subset_k) >= degree + 1) {
  model <- lm(y ~ poly(x, degree, raw = TRUE), data = subset_k)
  models[[k]] <- model
  fitted_data$y_hat[data$interval == k] <- predict(model, newdata = subset_k)
} else {
  models[[k]] <- NA
  fitted_data$y_hat[data$interval == k] <- NA
}
}

p <- ggplot(fitted_data, aes(x = x)) +
  geom_point(aes(y = y), color = "black", size = 1.5, alpha = 0.6) +
  geom_line(aes(y = y_hat), color = "red", size = 1.2) +
  labs(title = paste("Piecewise Polynomial Fit — K =", K, ", Degree =", degree),
       x = "x", y = "y") +
  theme_minimal()

return(p)
}

print(fit_piecewise_poly(data, K = 4, degree = 2))

```



In this piecewise polynomial regression model with $K=4$ intervals and polynomials of degree 2, the data range is divided into four equally sized sections over which separate quadratic models are fitted.

The resulting curve successfully captures local variations in the data, showing smooth curvature within each segment. The fit adapts well to the general trend and local fluctuations, thanks to the flexibility of quadratic polynomials.

However, since the polynomials are fitted independently in each segment, **there are discontinuities at the boundaries** between intervals. This is a common issue in piecewise models without continuity constraints.

Overall, the model balances flexibility and interpretability. Increasing K further or using higher-degree polynomials could improve the local fit, but also risks overfitting or creating even sharper jumps between segments.